

# Siren System Requisites

This document describes the functional and non-functional requirements for Siren.

## Architecture Overview




Figure 1: architecture

## 1. High-Level Requirements

### 1.1 Core Functionality

- **HL-F1** – Siren shall transparently intercept and analyze L2–L4 traffic on a designated Linux interface without requiring topology changes.
- **HL-F2** – Siren shall rely on an eBPF pipeline (XDP + TC + socket hooks) to observe, rewrite, mirror, or drop packets at line rate.
- **HL-F3** – Siren shall evaluate YAML rules and apply actions such as `drop`, `rewrite`, `delay`, `throttle`, `duplicate`, `inject`, and `log`.
- **HL-F4** – Siren shall allow user-defined manipulators written in Go to extend the built-in actions.
- **HL-F5** – Siren shall record intercepted traffic plus rule decisions in at least one structured format (PCAP or JSONL) for offline analysis.
- **HL-F6** – Siren shall optionally hook TLS/SSL functions to mirror plaintext buffers (keys or decrypted payloads) without handling private certificates.

### 1.2 Operator Experience

- **HL-U1** – Siren shall be configured through a single YAML *playbook* that covers targets, pipelines, recording, and rules.
- **HL-U2** – Siren shall run as a standalone CLI binary that requires `sudo` (or equivalent capabilities) to attach its eBPF programs.
- **HL-U3** – Siren shall expose an optional REST API for runtime inspection or hot-swapping of rules and recording settings.

### **1.3 System Qualities**

- **HL-Q1** – Siren shall keep per-packet latency overhead below  $10\ \mu\text{s}$  in userspace by delegating enforcement to the kernel whenever possible.
  - **HL-Q2** – Siren shall be delivered as a self-contained binary that only depends on a compatible Linux kernel ( $\geq 5.4$ ) and standard tooling.
  - **HL-Q3** – Potentially sensitive features (e.g., TLS interception, packet injection) shall be opt-in and disabled by default.
  - **HL-Q4** – Deployment shall avoid complex prerequisites such as `iptables` rewrites or `LD_PRELOAD` tricks; selecting an interface is sufficient.
- 

## **2. Low-Level Requirements**

### **2.1 General**

- **LL-G1** – The CLI shall accept `-config <path>` to load a playbook.
- **LL-G2** – The YAML schema shall be versioned and documented, covering eBPF attachment, TLS providers, targets, rules, manipulators, and recording.

### **2.2 eBPF Engine (SRN-EBPF)**

#### **2.2.1 Architectural Requirements**

- **SRN-EBPF-A1** – An XDP program shall provide read access to ingress packets and enforce immediate drops.
- **SRN-EBPF-A2** – A TC clsact program shall provide rewrite/inject capabilities on both ingress and egress directions.
- **SRN-EBPF-A3** – Socket/cgroup programs shall emit per-flow metadata (state, latency).
- **SRN-EBPF-A4** – A ring buffer (or similar channel) shall transfer packet metadata/payload slices to userspace.
- **SRN-EBPF-A5** – BPF maps shall receive enforcement instructions such as drop, rewrite window, or throttle budget.

#### **2.2.2 Functional Requirements**

- **SRN-EBPF-F1** – The XDP program shall filter packets based on a target map (IP, MAC, IP:Port, EtherCAT, labels). Empty maps imply “capture all.”
- **SRN-EBPF-F2** – TC hooks shall apply rewrite/inject directives referenced by ID, supplied by userspace.
- **SRN-EBPF-F3** – Socket hooks shall pace or mirror flows when instructed by the rule engine.
- **SRN-EBPF-F4** – Userspace shall manage the lifecycle of every loaded program (load, attach, detach, reload).

#### **2.2.3 Non-Functional Requirements**

- **SRN-EBPF-N1** – All programs shall support Linux kernels  $\geq 5.4$ .
- **SRN-EBPF-N2** – Precompiled objects shall be embedded into the Go binary.
- **SRN-EBPF-N3** – `go generate ./siren/ebpf` shall rebuild every object via `clang -target bpf`.

## 2.3 TLS Instrumentation (SRN-TLS)

- **SRN-TLS-A1** – Siren shall support optional uprobes/uretprobes on userland TLS libraries (OpenSSL, BoringSSL, etc.).
- **SRN-TLS-F1** – When enabled, plaintext buffers from `SSL_read/SSL_write` (or equivalent) shall be forwarded to the rule engine.
- **SRN-TLS-F2** – TLS providers shall be whitelisted in the YAML playbook by library path and symbol names.
- **SRN-TLS-N1** – If TLS instrumentation is disabled or a provider is missing, Siren shall continue operating on clear text flows normally and should skip ciphertext to avoid wasting resources analyzing encrypted data.

## 2.4 Intercept Engine (SRN-INT)

### 2.4.1 Architectural Requirements

- **SRN-INT-A1** – The intercept engine shall receive traffic contexts from the eBPF pipeline and enrich them with rule metadata.
- **SRN-INT-A2** – The engine shall evaluate matchers first, then hand off to manipulators before emitting the final action.

### 2.4.2 Functional Requirements

- **SRN-INT-F1** – Matchers shall include direction, payload regex, byte ranges, protocol/port, and custom tags.
- **SRN-INT-F2** – Actions shall include `drop`, `rewrite`, `delay`, `duplicate`, `mirror`, `throttle`, `inject`, and `log`.
- **SRN-INT-F3** – Actions backed by kernel enforcement (`drop/rewrite/throttle`) shall be reflected in the appropriate BPF maps; remaining actions shall run in userspace or manipulators.
- **SRN-INT-F4** – Rules shall be hot-swappable via the REST API without restarting the binary.

## 2.5 Recorder (SRN-REC)

### 2.5.1 Functional Requirements

- **SRN-REC-F1** – The recorder shall capture every intercepted packet to at least one destination (PCAP, JSONL).
- **SRN-REC-F2** – Recording shall be controlled through the YAML playbook and adjustable via the REST API.

### 2.5.2 Non-Functional Requirements

- **SRN-REC-N1** – Recording shall avoid blocking the interception loop; buffering or asynchronous writes shall be used to keep overhead negligible.